

Data 624: Week 4 Homework

Angrand, Burke, Deboch, Groysman, Karr

October 12, 2019

Week 4 Assignment

Chapter 7 HA 7.1, 7.3


7.1 Consider the *pigs* series - the number of pigs slaughtered in Victoria each month.

The `pigs` ts contains monthly total number of pigs slaughtered in Victoria, Australia (Jan 1980 to Aug 1995).

```
#sampling and shape of dataset, preliminary EDA
head(pigs)

##           Jan      Feb      Mar      Apr      May      Jun
## 1980  76378  71947  33873  96428 105084  95741

#autoplot the series
pigsdata<-window(pigs)
autoplot(pigsdata) +
  ylab("Number of pigs slaughtered") +xlab("Year")
```



a. Use the `ses()` function in R to find the optimal values of α and ℓ_0 , and generate forecasts for the next four months.

```
#Estimate parameters
fc_pigs_ses<-ses(pigsdata, h=4)

#Get forecasted estimate parameters from model -(alpha and L)
round(fc_pigs_ses$model$par[1:2],4)

##      alpha      1
##    0.2971 77260.0561

#generate 4 months of forecasts
data.frame(fc_pigs_ses)

##      Point.Forecast   Lo.80   Hi.80   Lo.95   Hi.95
## Sep 1995      98816.41 85605.43 112027.4 78611.97 119020.8
## Oct 1995      98816.41 85034.52 112598.3 77738.83 119894.0
## Nov 1995      98816.41 84486.34 113146.5 76900.46 120732.4
## Dec 1995      98816.41 83958.37 113674.4 76092.99 121539.8

#Accuracy of one-step-ahead training errors
round(accuracy(fc_pigs_ses),2)

##      ME      RMSE      MAE      MPE  MAPE  MASE  ACF1
## Training set 385.87 10253.6 7961.38 -0.92 9.27  0.8 0.01


# see how SES model was fitted
fc_pigs_ses$model

## Simple exponential smoothing
##
## Call:
## ses(y = pigsdata, h = 4)
##
## Smoothing parameters:
##   alpha = 0.2971
##
## Initial states:
##   l = 77260.0561
##
## sigma: 10308.58
##
##      AIC      AICc      BIC
## 4462.955 4463.086 4472.665

# get 1st 4 months of forecasts
tsCV(pigs,ses,h=4)[1:4,]
```

```
##           h=1           h=2           h=3           h=4
## [1,]   -4431.00  -42505.00  20050.00  28706.00
## [2,]  -42431.44   20123.56  28779.56  19436.56
## [3,]   62555.00   71211.00  61868.00  76774.00
## [4,]   28706.00   19363.00  34269.00  23953.00

#Plot (Note that forecast using ses doesn't have a trend component.)
fc_pigs_ses<-ses(pigs, h=100)
autoplot(fc_pigs_ses) +
  autolayer(fitted(fc_pigs_ses), series="fitted") +
  ylab("Total number of Pigs slaughtered in Victoria, Australia") + xlab("
Jan 1980 to Aug 1995")
```



b. Compute a 95% prediction interval for the first forecast using $\hat{y} \pm 1.96s$ is the standard deviation of the residuals. Compare your interval with the interval produced by R.

```
# 95% prediction interval for the first forecast

lower.upper <- data.frame( fc_pigs_ses$lower[1, "95%"], fc_pigs_ses$upper[1,
"95%"])
names(lower.upper)<- c("lower.limit", "upper.limit")
lower.upper

##      lower.limit upper.limit
## 95%      78611.97   119020.8
```

```

# calculate standard deviation of residuals with and without model,  $s = 10273.69$  vs  $s$  (estimated) 10308.58
s <- sd(fc_pigs_ses$residuals)
print(paste("Standard Deviation: ", round(s, 2)))

## [1] "Standard Deviation: 10273.69"

# calculate 95% prediction interval with model
pred.interval.model <- data.frame(fc_pigs_ses$mean[1] - 1.96*s, fc_pigs_ses$mean[1] + 1.96*s)
names(pred.interval.model) <- c("Lower.Model", "Upper.Model")
pred.interval.model

##      Lower.Model Upper.Model
## 1      78679.97    118952.8

# calculate 95% prediction interval without model
pred.interval.womodel <- data.frame(mean(pigs) - 1.96*s, mean(pigs) + 1.96*s)
names(pred.interval.womodel) <- c("Lower.NOModel", "Upper.NOModel")
pred.interval.womodel

##      Lower.NOModel Upper.NOModel
## 1           70504      110776.9

```

- R gives an interval of [78611.97, 119020.8] and by computing the standard deviation of the residuals we got [78679.97, 118952.8]. Both are really close.

7.3 Modify your function from the previous exercise to return the sum of squared errors rather than the forecast of the next observation. Then use the `optim()` function to find the optimal values of α and ℓ_0 . Do you get the same values as the `ses()` function?

- `my_ses_func` returns an ℓ value that is $\sim 0.01\%$ different than the `ses` function.

```

##      alpha      1
##      0.2971 77269.3253

##      alpha      1
##      0.2971 77260.0561

```