



PROJECT 1 FALL 2019

Group 1: Angrand, Burke, Deboch, Groysman, Karr

Table of Contents

Part A - ATM Forecast, ATM624Data.xlsx	2
Part B - Forecasting Power, ResidentialCustomerForecastLoad-624.xlsx	25
Part C – Waterflow_Pipe1.xlsx and Waterflow_Pipe2.xlsx	40
Step 1. Load Libraries.....	40
Step 2. Read in 2 Excel files.....	40
Step 3. Exploratory Analysis.....	40
Step 4. Data Cleaning.	44
Step 5. Converting data into time series.	46
Step 6. Looking at seasonality and trend.....	48
Step 8. Applying decomposition.	51
Step 9. Exponential Forecasting.....	55
Step 9. Selecting Forecasting Method.	57
Step 10. Preparing the final file to be ouputed in the Excel.....	60

Data 624. Project 1: Part A

Angrand, Burke, Deboch, Groysman, Karr

October 22, 2019

Part A - ATM Forecast, ATM624Data.xlsx

Data: ATM624Data.xlsx

In part A, I want you to forecast how much cash is taken out of 4 different ATM machines for May 2010. The data is given in a single file. The variable 'Cash' is provided in hundreds of dollars, other than that it is straight forward. I am being somewhat ambiguous on purpose. I am giving you data, please provide your written report on your findings, visuals, discussion and your R code all within a Word readable document, except the forecast which you will put in an Excel readable file. I must be able to cut and paste your R code and run it in R studio. Your report must be professional - most of all - readable, EASY to follow. Let me know what you are thinking, assumptions you are making! Your forecast is a simple CSV or Excel file that MATCHES the format of the data I provide.

```
#Upload Library
library(tidyverse)
library(readxl)
library(fpp2)
library(forecast)
```

Read in File/EDA/Data Adjustments

- drop nulls
- restructure dataset
- convert to time series object
- plot time series object for each ATM

```
temp = tempfile(fileext = ".xlsx")
dataURL <- "https://raw.githubusercontent.com/mburke65/CUNY_Data624/master/ProjectFolder/Provided_Files/ATM624Data.xlsx"
download.file(dataURL, destfile=temp, mode='wb')
```

```
atm <- readxl::read_excel(temp, sheet =1)
```

```
head(atm,5)
```

```
## # A tibble: 5 x 3
##   DATE                ATM    Cash
##   <dtm>              <chr> <dbl>
## 1 2009-05-01 00:00:00 ATM1     96
## 2 2009-05-01 00:00:00 ATM2    107
```

```

## 3 2009-05-02 00:00:00 ATM1      82
## 4 2009-05-02 00:00:00 ATM2      89
## 5 2009-05-03 00:00:00 ATM1      85

#Drop null values
atm<-atm %>%
  drop_na()

#Convert each ATM to Column
atm<- atm %>%
  spread(ATM,Cash)
head(atm,5)

## # A tibble: 5 x 5
##   DATE                ATM1  ATM2  ATM3  ATM4
##   <dtm>              <dbl> <dbl> <dbl> <dbl>
## 1 2009-05-01 00:00:00    96   107     0    96
## 2 2009-05-02 00:00:00    82    89     0    82
## 3 2009-05-03 00:00:00    85    90     0    85
## 4 2009-05-04 00:00:00    90    55     0    90
## 5 2009-05-05 00:00:00    99    79     0    99

#Fix the date column
atm <- atm %>%
  mutate(DATE =as.Date(DATE))
head(atm)

## # A tibble: 6 x 5
##   DATE                ATM1  ATM2  ATM3  ATM4
##   <date>              <dbl> <dbl> <dbl> <dbl>
## 1 2009-05-01         96   107     0    96
## 2 2009-05-02         82    89     0    82
## 3 2009-05-03         85    90     0    85
## 4 2009-05-04         90    55     0    90
## 5 2009-05-05         99    79     0    99
## 6 2009-05-06         88    19     0    88

#Convert to a time series
ts_atm <- ts(atm %>% select(-DATE))

head(ts_atm)

## Time Series:
## Start = 1
## End = 6
## Frequency = 1
##   ATM1 ATM2 ATM3 ATM4
## 1   96  107    0   96

```

```
## 2    82    89    0    82
## 3    85    90    0    85
## 4    90    55    0    90
## 5    99    79    0    99
## 6    88    19    0    88
```

- ATM1, ATM2, and ATM4 are a big deal of variation. but ATM3 shows no cash withdrawn for most of the year. One assumption we can do about ATM3 is that it has just opened. we will use the entire time series of ATM1 and ATM2. ATM3 will be used to forecast future prediction.

#Separate each ATM from the dataset and graph each dataset

```
atm1<-ts_atm[, "ATM1"]
```


```
autoplot(atm1) +
```

```
  labs(title = "Cash Withdrawn from ATM1", x="Day") +
```

```
  scale_y_continuous("Cash withdrawals in hundreds") +
```

```
  scale_color_discrete(NULL)
```

Cash Withdrawn from ATM1



#Separate each ATM from the dataset and graph each dataset


```
atm2<-ts_atm[, "ATM2"]
```

```
autoplot(atm2) +
```

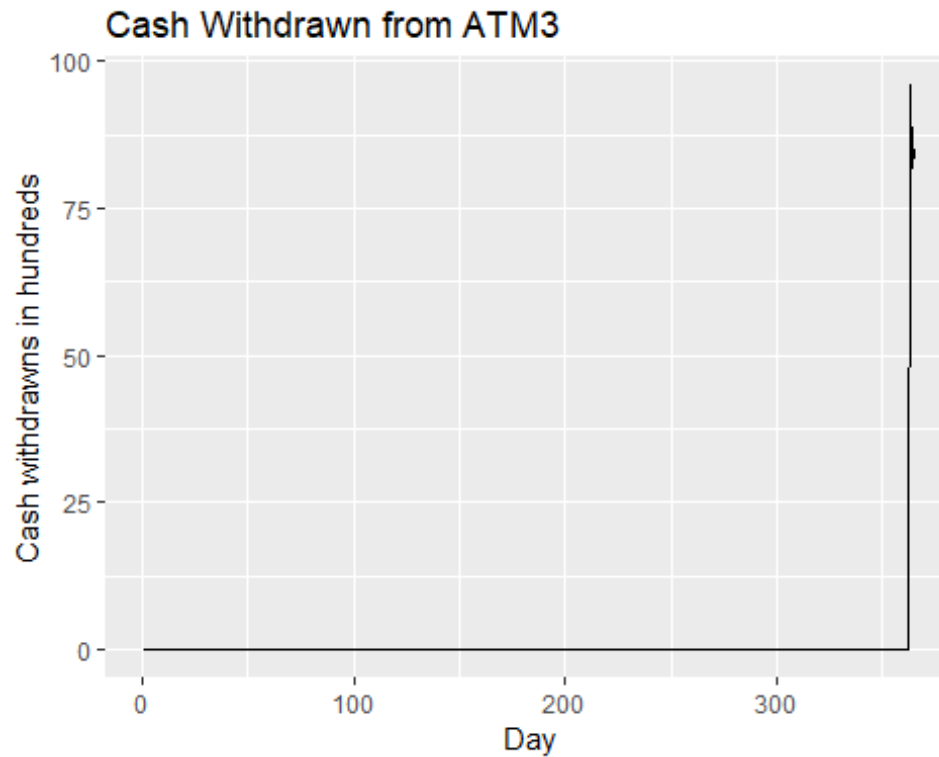
```
  labs(title = "Cash Withdrawn from ATM2", x="Day") +
```

```
  scale_y_continuous("Cash withdrawals in hundreds") +
```


```
  scale_color_discrete(NULL)
```



```
#Separate each ATM from the dataset and graph each dataset  
atm3<-ts_atm[,"ATM3"]  
autoplot(atm3) +  
  labs(title ="Cash Withdrawn from ATM3", x="Day") +  
  scale_y_continuous("Cash withdrawals in hundreds") +  
  scale_color_discrete(NULL)
```




```
#Separate each ATM from the dataset and graph each dataset  
atm4<-ts_atm[,"ATM4"]  
autoplot(atm4) +  
  labs(title = "Cash Withdrawn from ATM4", x="Day") +  
  scale_y_continuous("Cash withdrawals in hundreds") +  
  scale_color_discrete(NULL)
```



- ATM1, ATM2 and ATM4 show a lot of deal of seasonality in the withdrawn from those ATM. We can further analyze it by selecting the first 2 months of the data.

```
autoplot(ts(ts_atm[1:61, ])) +  
  
  labs(title = "Cash withdrawn from 4 ATMs",  
        subtitle = "May 2009 - June 2010",  
        x = "Day") +  
  
  scale_y_continuous("Cash withdrawn (hundreds)") +  
  
  scale_color_discrete(NULL)
```

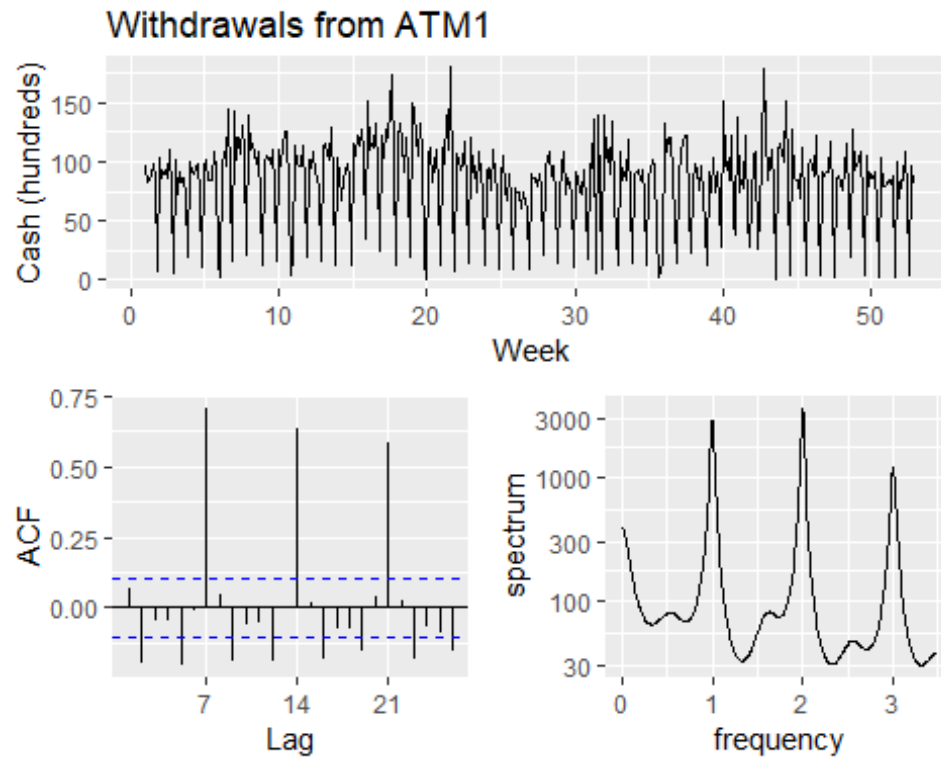
- The data presents a sort of weekly seasonality. To capture the seasonality of this data we will set the frequency to 7.

```
atm1_freq<-ts(atm1, frequency =7)
atm2_freq<-ts(atm2, frequency=7)
atm4_freq<-ts(atm4, frequency=7)
```

ATM1

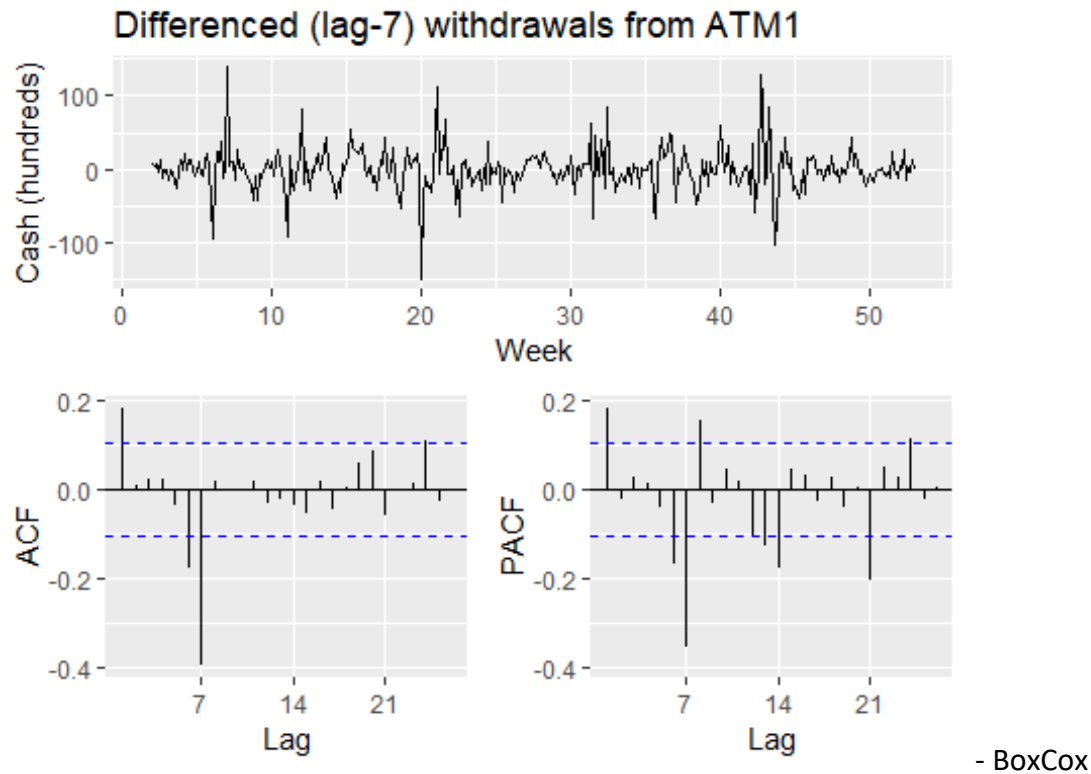
#ACF and spectrum plot

```
ggtsdisplay(atm1_freq, points = FALSE, plot.type = "spectrum",
            main = "Withdrawals from ATM1", xlab = "Week", ylab = "Cash (hundreds)")
```



- In 7, 14 and 21 there are large spikes. the frequency 1,2,3 show the spike as well. Both suggest a seasonal ARIMA model.

```
ggtsdisplay(diff(atm1_freq, 7), points = FALSE,
             main = "Differenced (lag-7) withdrawals from ATM1",
             xlab = "Week", ylab = "Cash (hundreds)")
```



transformation to estimate lambda

```
# get optimal Lambda for Box-cox transformation
lambda_atm1<- BoxCox.lambda(atm1_freq)

# define function to create models & return AIC values for timeseries
aic_atm<- function(p, d, q, P, D, Q) {

  # create model with Box-Cox and specified ARIMA parameters; extract AIC

  AIC(Arima(atm1_freq, order = c(p, d, q), seasonal = c(P, D, Q), lambda = lambda_atm1))

}

# create possible combinations of p, q, P, Q except all zero
expand.grid(p = 0:1, q = 0:1, P = 0:1, Q = 0:1) %>%

  filter(p > 0 | q > 0 | P > 0 | Q > 0) %>%

  # calc AIC for models
```

```

mutate(aic = pmap_dbl(list(p, 0, q, P, 1, Q), aic_atm)) %>%
# return best AIC

slice(which.min(aic))

##   p q P Q   aic
## 1 1 1 0 1 1221.26

```

- The minimum aic value is for non-seasonality AR(1) and MA(1). AR(0) and AM(1) is for seasonality. Let's fit the model using arima model arima(1,0,1)(0,1,1)

```

fit_atm1 <- Arima(atm1_freq, order = c(1, 0, 1), seasonal = c(0, 1, 1)
, lambda = lambda_atm1)
summary(fit_atm1)

## Series: atm1_freq
## ARIMA(1,0,1)(0,1,1)[7]
## Box Cox transformation: lambda= 0.2584338
##
## Coefficients:
##           ar1          ma1          sma1
##      -0.4894    0.6125   -0.6385
## s.e.    0.2309    0.2081    0.0432
##
## sigma^2 estimated as 1.732:  log likelihood=-606.63
## AIC=1221.26   AICc=1221.37   BIC=1236.78
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set 2.293003 24.81988 15.66437 -89.57546 108.1682 0.892827
##              ACF1
## Training set -0.008839946

```

Let's diagnostic the residuals with Ljung-Box.

```

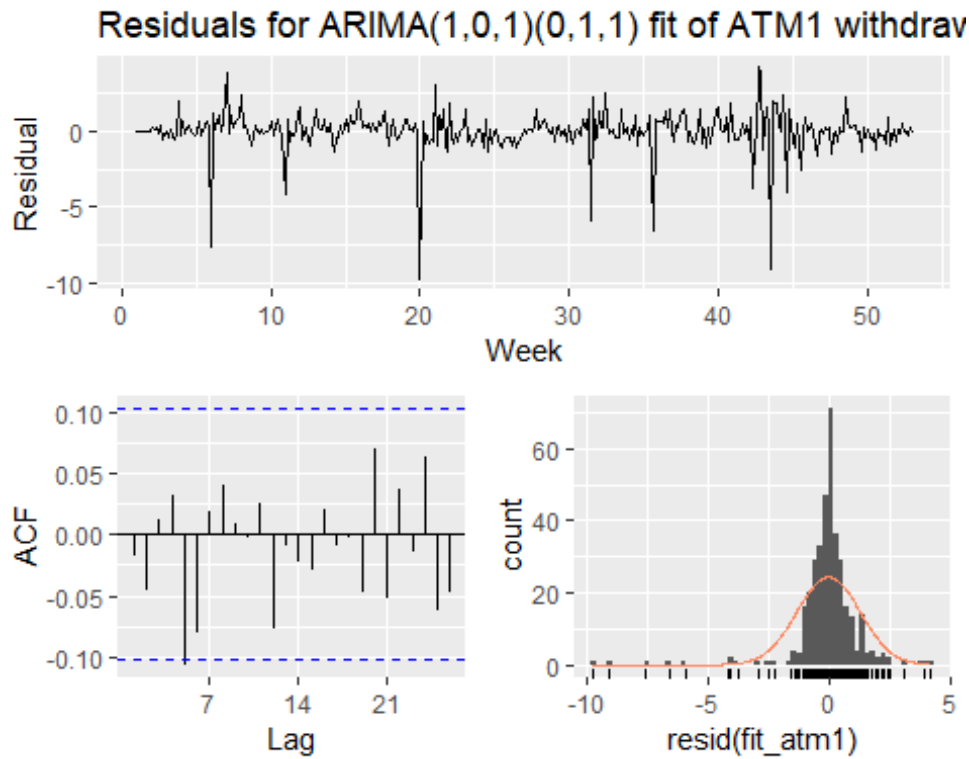
Box.test(resid(fit_atm1), type = "L", fitdf = 3, lag = 7)

##
## Box-Ljung test
##
## data:  resid(fit_atm1)
## X-squared = 8.0497, df = 4, p-value = 0.08977

ggtsdisplay(resid(fit_atm1), points = FALSE, plot.type = "histogram",
            main = "Residuals for ARIMA(1,0,1)(0,1,1) fit of ATM1 with
drawals",


```

```
xlab = "Week", ylab = "Residual")
```



- The p -value is greater than 0.05 meaning that the residual is white noise. The residuals are not correlated and there is a normal distribution around the mean 0. We can use that model for forecasting.

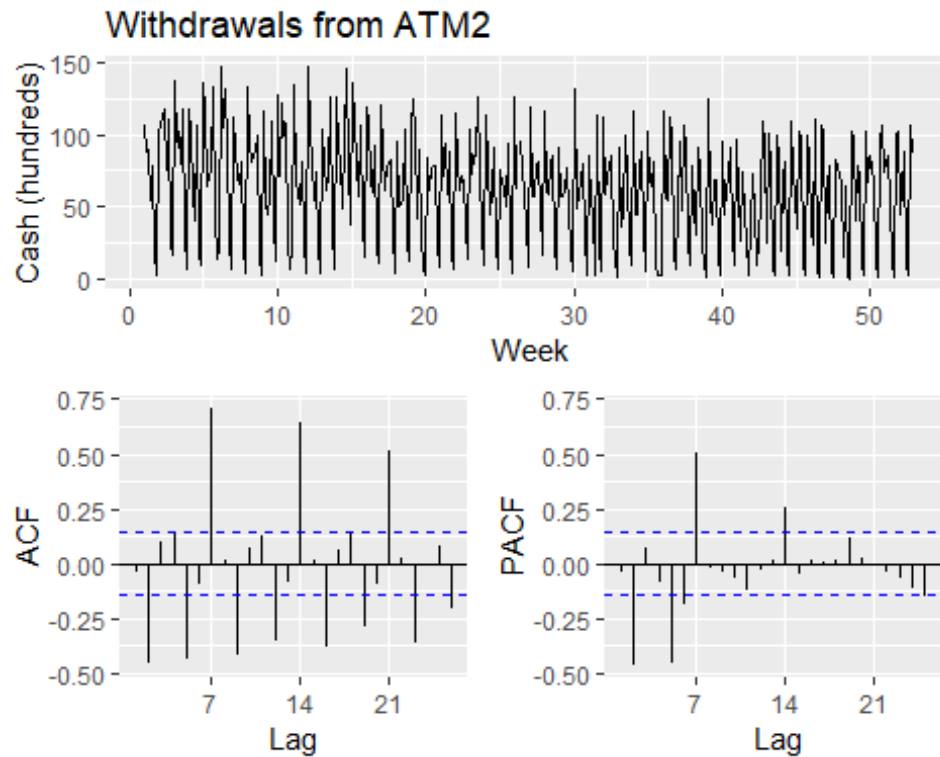
```
forecast_atm1 <- forecast(fit_atm1, 31, level = 95)
autoplot(forecast_atm1) +
  labs(title = "ATM1: ARIMA(1,0,1)(0,1,1)", x = "Week", y = NULL) +
  theme(legend.position = "none")
```



ATM2


- We can repeat the same step for ATM2.

```
ggtsdisplay(atm2_freq, points = FALSE,  
            main = "Withdrawals from ATM2", xlab = "Week", ylab = "Cash (hundreds)")
```



- The lag difference is 7.

```
ggtsdisplay(diff(atm2_freq, 7), points = FALSE,  
             main = "Differenced (lag-7) withdrawals from ATM2",  
             xlab = "Week", ylab = "Cash (hundreds)")
```



- The spikes in ACF & PACF in the non-differenced series at $k = 2$ & $k = 5$ suggest $p, q \in [0, 2, 5]$. using the same aic function we can evaluate the minimum aic

get optimal lambda for Box-cox transformation

```
lambda_atm2 <- BoxCox.lambda(atm2_freq)
```

Evaluate aic

```
aic_atm <- function(p, d, q, P, D, Q) {
```

```
  # create model with Box-Cox and specified ARIMA parameters; extract AIC
```

```
  AIC(Arima(atm2_freq, order = c(p, d, q), seasonal = c(P, D, Q), lambda = lambda_atm2))
```

```
}
```

create possible combinations of p, q, P, Q except all zero

```
expand.grid(p = c(0, 2, 5), q = c(0, 2, 5), P = 0:1, Q = 0:1) %>%
```

```
  filter(p > 0 | q > 0 | P > 0 | Q > 0) %>%
```

```
  # calculate AIC for models
```



```

mutate(aic = pmap_dbl(list(p, 0, q, P, 1, Q), aic_atm)) %>%
# return minimum AIC

slice(which.min(aic))

##    p q P Q      aic
## 1 2 2 0 1 2323.517

```

- the model arima used is arima(5,0,5)(0,1,1). Let's fit that model

```

fit_atm2<-Arima(atm2_freq, order = c(5, 0, 5), seasonal = c(0, 1, 1),
lambda = lambda_atm2)
summary(fit_atm2)

## Series: atm2_freq
## ARIMA(5,0,5)(0,1,1)[7]
## Box Cox transformation: lambda= 0.6584081
##
## Coefficients:
##          ar1          ar2          ar3          ar4          ar5          ma1          ma2
##      ma3
##          0.2055   -0.1209   0.2260   0.3032   -0.4312   -0.1448   0.0114   -0
##      .2213
## s.e.    0.4529    0.4033    0.2176    0.2419    0.4136    0.4787    0.4200    0
##      .2100
##          ma4          ma5          sma1
##          -0.2466    0.2470   -0.6905
## s.e.    0.2463    0.4176    0.0595
##
## sigma^2 estimated as 37.91:  log likelihood=-1152.1
## AIC=2328.19   AICc=2329.1   BIC=2374.76
##
## Training set error measures:
##              ME          RMSE          MAE   MPE  MAPE          MASE
## ACF1
## Training set 0.2238867 23.87153 16.65107 -Inf  Inf 0.8279025 -0.030
## 50682

```

- Let's evaluate the residual to check the validity of the model

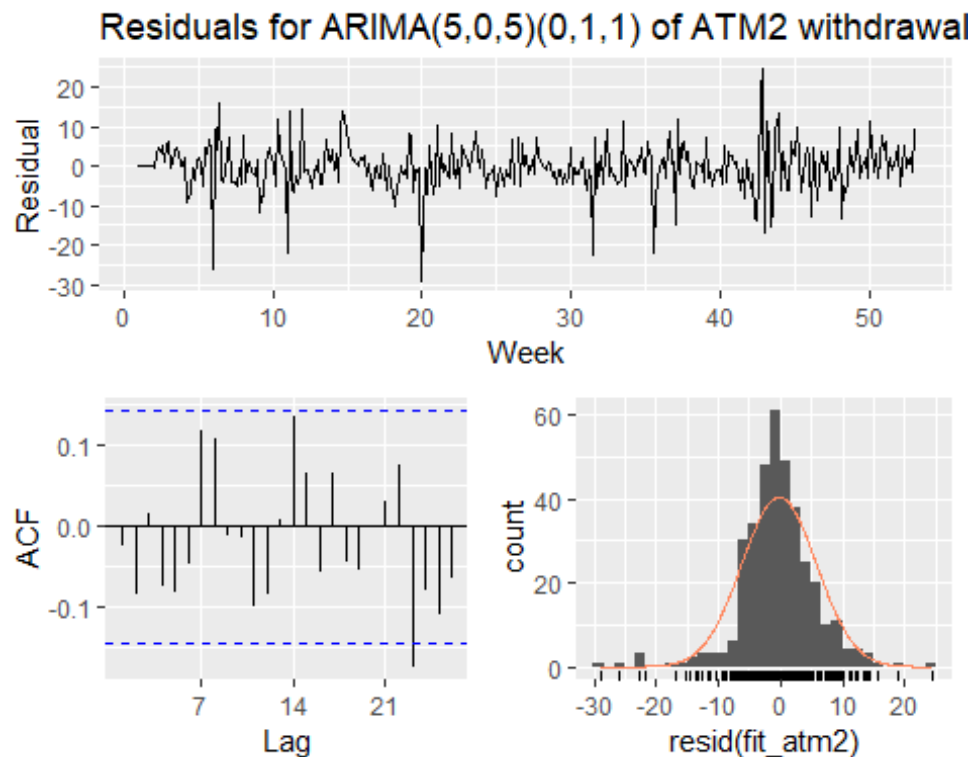
```

Box.test(resid(fit_atm2), type = "L", fitdf = 11, lag = 14)

##
## Box-Ljung test
##
## data:  resid(fit_atm2)
## X-squared = 2.1119, df = 3, p-value = 0.5495

```


```
ggtsdisplay(resid(fit_atm2), points = FALSE, plot.type = "histogram",
            main = "Residuals for ARIMA(5,0,5)(0,1,1) of ATM2 withdrawal",
            xlab = "Week", ylab = "Residual")
## Warning: Removed 1 rows containing non-finite values (stat_bin).
```



- P-value is greater than 0.05 and the residual appear to be normally distributed with a mean of 0. It can be used for forecast ATM2.

```
forecast_atm2<- forecast(fit_atm2, 31, level = 95)
autoplot(forecast_atm2) +
  labs(title = "ATM2: ARIMA(5,0,5)(0,1,1)", x = "Week", y = NULL) +
  theme(legend.position = "none")
```

ATM2: ARIMA(5,0,5)(0,1,1)



ATM4


- ATM4 has the same seasonality as ATM1 and ATM2. We will use the previous step to evaluate ATM4 model.

#Minimize the effect of the big withdraw in the day by using the median of the ATM4 dataset

```
atm4_freq[which.max(atm4_freq)] <- median(atm4_freq, na.rm = TRUE)
```

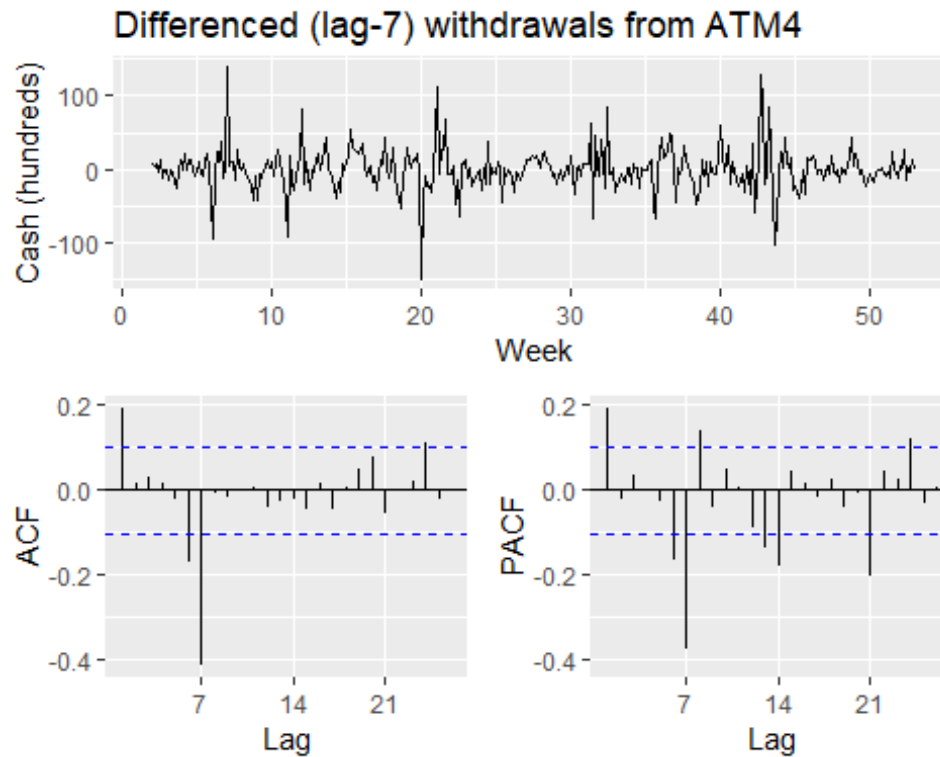
```
ggtsdisplay(atm4_freq, points = FALSE,
```

```
            main = "Withdrawals from ATM4", xlab = "Week", ylab = "Cash (hundreds)")
```



- We notice a difference lag of 7.

```
ggtsdisplay(diff(atm4_freq, 7), points = FALSE,  
             main = "Differenced (lag-7) withdrawals from ATM4",  
             xlab = "Week", ylab = "Cash (hundreds)")
```



- ARIMA model for ATM4 will be evaluated.

```
# get optimal lambda for Box-cox transformation
```

```
lambda_atm4 <- BoxCox.lambda(atm4_freq)
```

```
aic_atm(0,2,5,0,2,5)
```

```
## [1] 2365.837
```

```
# create possible combinations of p, q, P, Q except all zero
```

```
expand.grid(p = c(0, 2, 5), q = c(0, 2, 5), P = 0:1, Q = 0:1) %>%
```

```
  filter(p > 0 | q > 0 | P > 0 | Q > 0) %>%
```

```
# calculate AIC for models
```

```
mutate(aic = pmap_dbl(list(p, 0, q, P, 1, Q), aic_atm)) %>%
```

```
# return minimum AIC
```

```
slice(which.min(aic))
```

```
##   p q P Q      aic
```

```
## 1 2 2 0 1 2323.517
```

- Let's fit the ARIMA model with the values (0,0,2)(0,1,1)

```
fit_atm4<-Arima(atm4_freq, order = c(0, 0, 2), seasonal = c(0, 1, 1),
lambda = lambda_atm4)
summary(fit_atm4)

## Series: atm4_freq
## ARIMA(0,0,2)(0,1,1)[7]
## Box Cox transformation: lambda= 0.2355973
##
## Coefficients:
##          ma1          ma2          sma1
##          0.1094      -0.1089      -0.6468
## s.e.    0.0524      0.0523      0.0422
##
## sigma^2 estimated as 1.467:  log likelihood=-576.96
## AIC=1161.92   AICc=1162.03   BIC=1177.44
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MAS
## Training set 2.356651 24.88094 15.90136 -85.71176 104.5953 0.902312
##
##              ACF1
## Training set 0.02127326
```


- Let's investigate the residuals using Ljung-box test

```
Box.test(resid(fit_atm4), type = "L", fitdf = 3, lag = 7)

##
## Box-Ljung test
##
## data:  resid(fit_atm4)
## X-squared = 5.7899, df = 4, p-value = 0.2154

ggtsdisplay(resid(fit_atm4), points = FALSE, plot.type = "histogram",
            main = "Residuals for ARIMA(0,0,2)(0,1,1) of ATM4 withdrawals",
            xlab = "Week", ylab = "Residual")
```

Residuals for ARIMA(0,0,2)(0,1,1) of ATM4 withdrawals



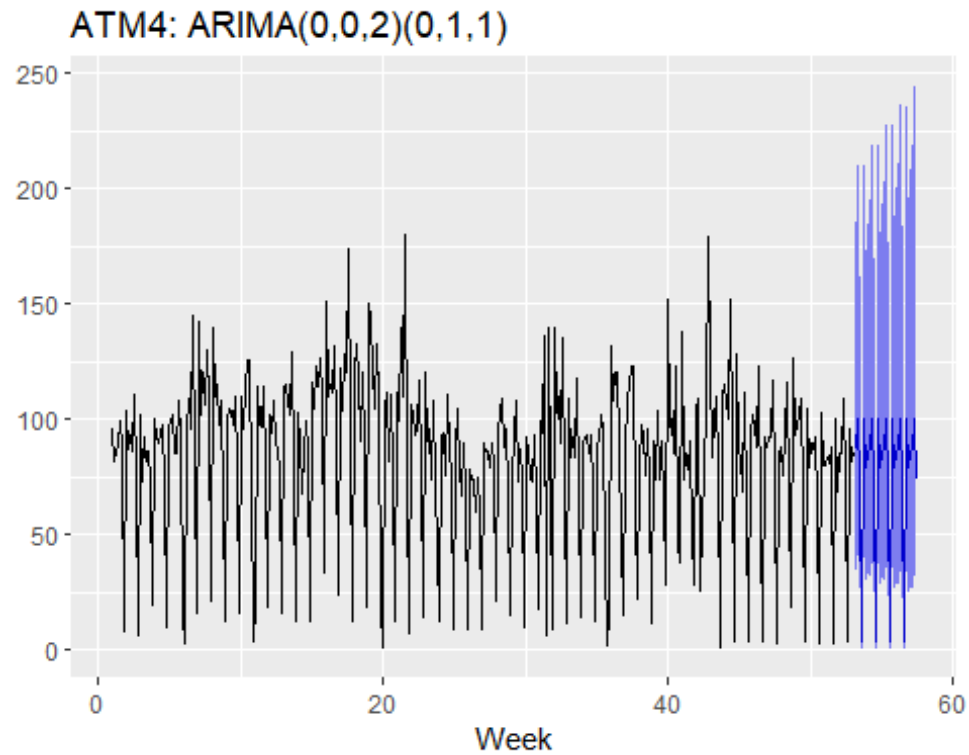
- It is normally distributed around a mean of 0. p-value is also greater than 0.05. We can use the model to forecast.

```
forecast_atm4<- forecast(fit_atm4, 31, level = 95)
```

```
autoplot(forecast_atm4) +
```

```
  labs(title = "ATM4: ARIMA(0,0,2)(0,1,1)", x = "Week", y = NULL) +
```

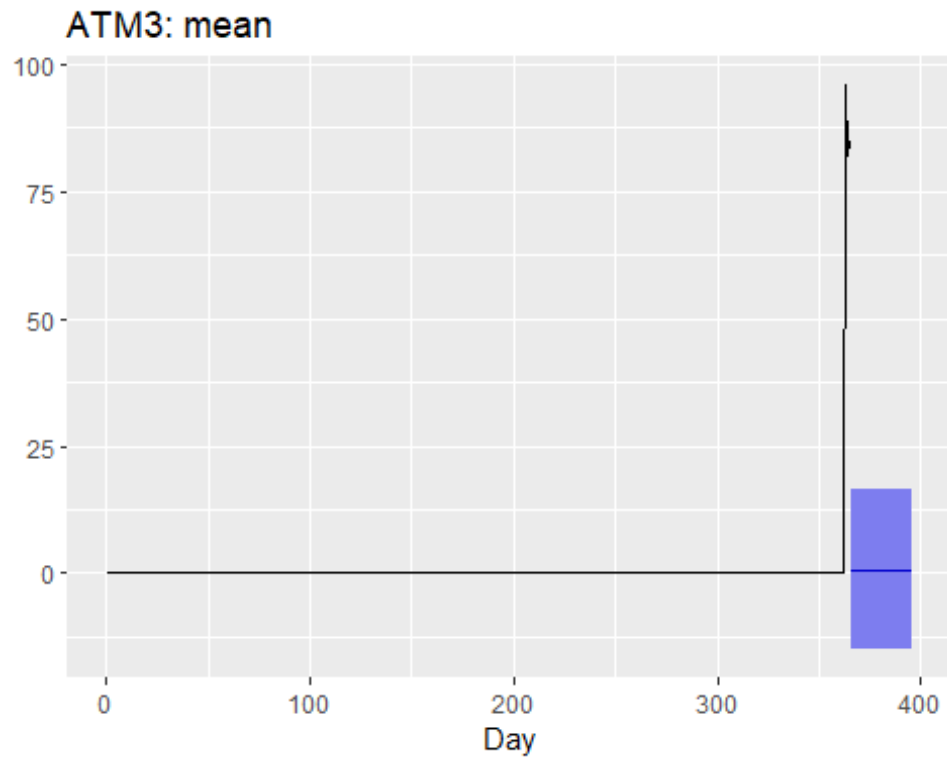
```
  theme(legend.position = "none")
```



ATM3

- Since ATM3 contains limited data we will use the mean forecast method.

```
forecast_atm3 <- meanf(atm3, 31, level = 95)
autoplot(forecast_atm3) +
  labs(title = "ATM3: mean", x = "Day", y = NULL) +
  theme(legend.position = "none")
```

Writing the forecast to a CSV file

```
data_frame(DATE = rep(max(atm$DATE) + 1:31, 4),
            atm = rep(names(atm)[-1], each = 31),
            Cash = c(forecast_atm1$mean, forecast_atm2$mean,
                    forecast_atm3$mean, forecast_atm4$mean)) %>%
write_csv("project1_forecast_atm_1.csv")
```

Warning: `data_frame()` is deprecated, use `tibble()`.

This warning is displayed once per session.

Data 624. Project 1: Part B

Angrand, Burke, Deboch, Groysman, Karr

October 22, 2019

Part B - Forecasting Power, ResidentialCustomerForecastLoad-624.xlsx

Data: ResidentialCustomerForecastLoad-624.xlsx

Part B consists of a simple dataset of residential power usage for January 1998 until December 2013. Your assignment is to model these data and a monthly forecast for 2014. The data is given in a single file. The variable 'KWH' is power consumption in Kilowatt hours, the rest is straight forward. Add these to your existing files above - clearly labeled.

```
library(httr)
library(xlsx)
library(ggplot2)
library(gridExtra)
library(forecast)
```

a. Read in File

```
temp = tempfile(fileext = ".xlsx")
dataURL <- "https://raw.githubusercontent.com/mburke65/CUNY_Data624/master/ProjectFolder/Provided_Files/ResidentialCustomerForecastLoad-624.xlsx"
download.file(dataURL, destfile=temp, mode='wb')

power.data <- readxl::read_excel(temp, sheet =1)
```

b. EDA Analysis

- Check/Fill in null values
- Convert to time series
- Graph the monthly data
 - General plot & seasonal plot: seasonality can be observed in the below plot. There are spikes each year from May to August (air conditioning?) and again in December (holiday season?). There is a slight dip in Jul 2010 maybe due to an unseasonably cold month.
 - Seasonal Box Plot: provides a similar visual to the seasonal plot with usage spikes in the summer months and December. IT also highlights the fluctuations in consumption within each month.
 - Decomposition components graph: this plot again shows that there is a general upwards trend in the data with an observed outlier in July 2010.

```
head(power.data)
```

```
## # A tibble: 6 x 3
##   CaseSequence `YYYY-MMM`      KWH
##   <dbl> <chr>      <dbl>
## 1      733 1998-Jan    6862583
## 2      734 1998-Feb    5838198
## 3      735 1998-Mar    5420658
## 4      736 1998-Apr    5010364
## 5      737 1998-May    4665377
## 6      738 1998-Jun    6467147
```

```
summary(power.data)
```

```
##   CaseSequence      YYYY-MMM      KWH
##   Min.   :733.0   Length:192   Min.    : 770523
##   1st Qu.:780.8   Class :character 1st Qu.: 5429912
##   Median :828.5   Mode  :character Median : 6283324
##   Mean   :828.5                Mean   : 6502475
##   3rd Qu.:876.2                3rd Qu.: 7620524
##   Max.   :924.0                Max.    :10655730
##                                     NA's    :1
```

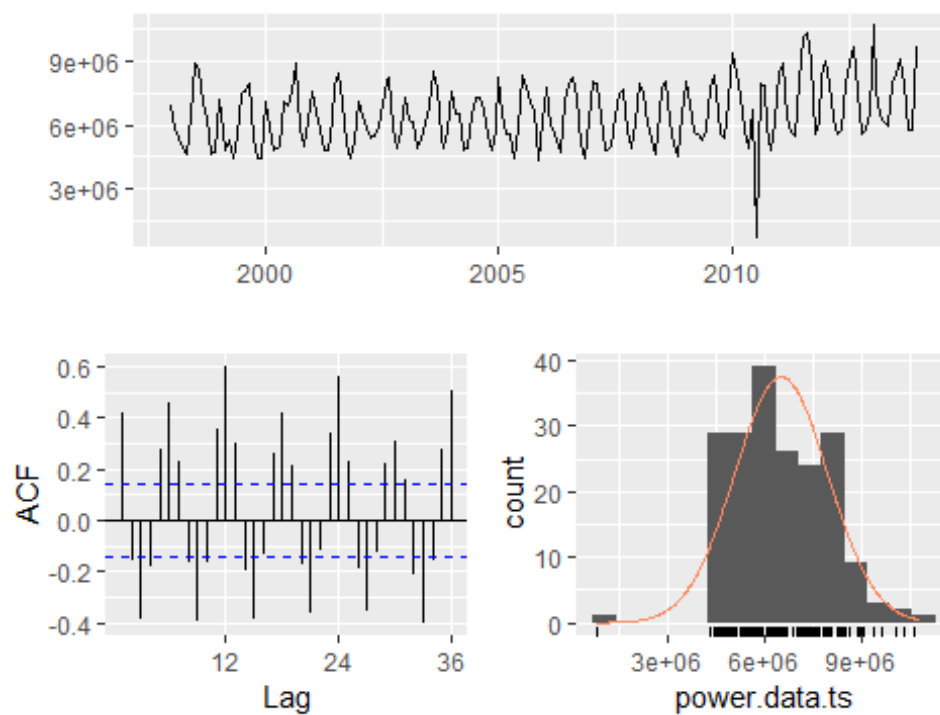
```
print(paste("Check for nulls: ",sum(is.na(power.data)), " Row of Nulls
"))
```

```
## [1] "Check for nulls: 1 Row of Nulls"
```


```
power.data[is.na(power.data)] <- median(power.data$KWH,na.rm = TRUE)
```

```
power.data.ts <-ts(power.data[, "KWH"],start = c(1998,1),frequency = 12
)
```


```
ggtsdisplay(power.data.ts, points = FALSE, plot.type = "histogram")
```



```
ggseasonplot(power.data.ts)
```




```
boxplot(power.data.ts~cycle(power.data.ts),xlab="Month", ylab = "Power Usage")
```



```
plot(decompose(power.data.ts))
```

Decomposition of additive time series



C. Model 1: Arima W/ Box-Cox Transformation

- Replace outlier with `tsoutlier` suggestion (utilizes a box-cox transformation)
- Use an auto arima model on the box-cox adjusted data
 - Suggested model: `ARIMA(0,0,3)(2,1,0)[12]` with drift. RSME(595389) & AICc (5332.67)
- Check the residuals to make sure the model is satisfactory:
 - ACF /PACF Plots: the residual appears normal residuals mostly around 0, suggesting stationarity of the residuals
 - The Box Ljung tests presents a p-value of 0.6951 which indicates white noise
- Forecast 2014 power values & plot forecasted values

#outlier detection/suggestion/replacement

```
find.outlier<- tsoutliers(power.data.ts, iterate = 2, lambda = "auto")
power.data.ts.bc<- power.data.ts
power.data.ts.bc[find.outlier$`index`[1]] <- find.outlier$replacements[1]
print(paste("Suggested/Implemented Change for Outlier: ",power.data.ts
.bc[151], " Original Value",power.data.ts[151]))

## [1] "Suggested/Implemented Change for Outlier: 7757388.48810024 0
riginal Value 770523"
```

```

#auto arima model
power.model <- auto.arima(power.data.ts.bc, seasonal = TRUE, stepwise
= FALSE)
summary.arima<- summary(power.model)


## Series: power.data.ts.bc
## ARIMA(0,0,3)(2,1,0)[12] with drift
##
## Coefficients:
##          ma1          ma2          ma3          sar1          sar2          drift
##          0.3492  0.0587  0.2303  -0.7222  -0.4251  9027.233
## s.e.  0.0788  0.0892  0.0741  0.0765  0.0784  3057.838
##
## sigma^2 estimated as 3.912e+11:  log likelihood=-2659.01
## AIC=5332.02  AICc=5332.67  BIC=5354.37
##
## Training set error measures:
##              ME    RMSE      MAE      MPE      MAPE      MAS
E
## Training set -8181.906 595389 434520.9 -0.8060555 6.610412 0.697749
1
##              ACF1
## Training set -0.01026567

summary.arima

##              ME    RMSE      MAE      MPE      MAPE      MAS
E
## Training set -8181.906 595389 434520.9 -0.8060555 6.610412 0.697749
1
##              ACF1
## Training set -0.01026567

#check residuals
checkresiduals(power.model)

```




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,0,3)(2,1,0)[12] with drift
## Q* = 14.513, df = 18, p-value = 0.6951
##
## Model df: 6.    Total lags used: 24

#forecast model @ 95%
forecast.power <- forecast(power.model, level = c(95), h =12)
forecast.power
```

	Point Forecast	Lo 95	Hi 95
## Jan 2014	10312755	9086940	11538570
## Feb 2014	8685531	7387130	9983933
## Mar 2014	7203085	5902687	8503482
## Apr 2014	6000251	4669575	7330927
## May 2014	5941905	4611229	7272581
## Jun 2014	8204931	6874255	9535607
## Jul 2014	9501418	8170742	10832094
## Aug 2014	9992966	8662290	11323642
## Sep 2014	8493959	7163283	9824635
## Oct 2014	5871672	4540996	7202348
## Nov 2014	6154352	4823676	7485028
## Dec 2014	8381806	7051130	9712482


```
autoplot(forecast.power)
```



D. Model 2: ETS W/ Box-Cox Transformation

- The ets function automatically selects the best method for forecasting data. the ets function selected ETS(M,N,M) exponential smoothing:
 - The first letter denotes the error type: multiplicative errors
 - The second letter denotes the trend type: no trend
 - The third letter denotes the season type: multiplicative seasonality
- Utilize the transformed data & ETS model
- Model Results: RMSE (630869.7) & AICc (6148.032)
- Check the residuals to make sure the model is satisfactory:
 - ACF /PACF Plots: the residual appears normal residuals mostly around 0, suggesting stationarity of the residuals
 - The Box Ljung tests presents a p-value of 0.0002921 which may indicate that there's dependency issues with the lags

#model w/ previously transformed data

```
power.model.ets <- ets(power.data.ts.bc)
summary.ets<- summary(power.model.ets)
```

```
## ETS(M,N,M)
```

```
##
```

```
## Call:
```

```

## ets(y = power.data.ts.bc)
##
## Smoothing parameters:
##   alpha = 0.1206
##   gamma = 0.203
##
## Initial states:
##   l = 6188160.6435
##   s = 0.9017 0.755 0.9295 1.223 1.2676 1.2298
##       1.0165 0.7614 0.8029 0.8903 1.029 1.1935
##
## sigma: 0.0971
##
##      AIC      AICc      BIC
## 6145.305 6148.032 6194.167
##
## Training set error measures:
##
##           ME      RMSE      MAE      MPE      MAPE      MA
SE
## Training set 51458.11 630869.7 482886.1 0.04287257 7.292037 0.77541
35
##           ACF1
## Training set 0.2096574


summary.ets

##           ME      RMSE      MAE      MPE      MAPE      MA
SE
## Training set 51458.11 630869.7 482886.1 0.04287257 7.292037 0.77541
35
##           ACF1
## Training set 0.2096574

#check residuals
checkresiduals(power.model.ets)

```

Residuals from ETS(M,N,M)

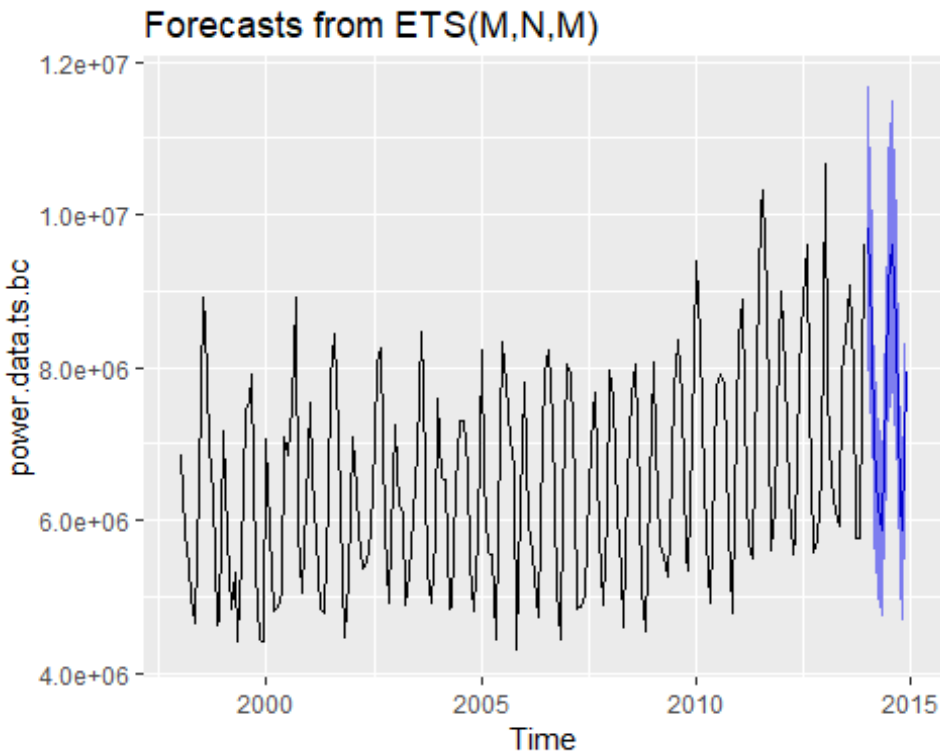


```
##
##  Ljung-Box test
##
## data:  Residuals from ETS(M,N,M)
## Q* = 32.819, df = 10, p-value = 0.0002921
##
## Model df: 14.    Total lags used: 24

#forecast model @ 95%
forecast.power.ets <- forecast(power.model.ets, level = c(95), h =12)
forecast.power.ets
```

```
##          Point Forecast    Lo 95    Hi 95
## Jan 2014      9825114 7955694 11694534
## Feb 2014      8460361 6838842 10081879
## Mar 2014      6974291 5627960  8320623
## Apr 2014      6167737 4968643  7366830
## May 2014      5886368 4733958  7038779
## Jun 2014      7783200 6248904  9317496
## Jul 2014      9070884 7270556 10871212
## Aug 2014      9599368 7681334 11517403
## Sep 2014      8501578 6791613 10211542
## Oct 2014      6241977 4978271  7505684
## Nov 2014      5885873 4686553  7085194
## Dec 2014      7933193 6306380  9560005
```

```
autoplot(forecast.power.ets)
```



E. Model 3: STLF

- STLF model will be the third model as it provides the user more control and can be robust when dealing with outliers. the STLF utilizes a local weighted regression to fit the points (Loess smoothing) and forecast future values.
- Model summary: RMSE (843670.1) & AICc (6255.445)
- Check residuals:
 - ACF/PACF: most lags are within the error bounds, suggesting stationarity of the residuals
 - Box Ljung:p-value = 0.1457 which indicates white noise
- Forecast 2014 power values & plot forecasted values

```
power.model.stl <- stlf(power.data.ts, s.window='periodic', robust=
TRUE)
summary.stl<- summary(power.model.stl)

##
## Forecast method: STL + ETS(A,N,N)
##
## Model Information:
## ETS(A,N,N)
##
```

```

## Call:
## ets(y = x, model = etsmodel, allow.multiplicative.trend = allow.mu
ltiplicative.trend)
##
## Smoothing parameters:
## alpha = 0.0892
##
## Initial states:
## l = 6317161.2015
##
## sigma: 848098.8
##
## AIC AICc BIC
## 6255.318 6255.445 6265.090
##
## Error measures:
## ME RMSE MAE MPE MAPE MAS
E
## Training set 69834.05 843670.1 512067.7 -4.243142 12.03155 0.731642
2
## ACF1
## Training set 0.209786
##
## Forecasts:
## Point Forecast Lo 80 Hi 80 Lo 95 Hi 95
## Jan 2014 8919230 7832347 10006112 7256987 10581473
## Feb 2014 7833393 6742199 8924586 6164556 9502230
## Mar 2014 7005154 5909665 8100642 5329749 8680558
## Apr 2014 6358706 5258940 7458473 4676759 8040654
## May 2014 6086344 4982317 7190371 4397880 7774808
## Jun 2014 7653295 6545023 8761567 5958339 9348251
## Jul 2014 8801193 7688692 9913693 7099770 10502616
## Aug 2014 9301580 8184867 10418293 7593714 11009445
## Sep 2014 8524478 7403568 9645387 6810194 10238761
## Oct 2014 6593256 5468165 7718347 4872577 8313934
## Nov 2014 5961092 4831835 7090349 4234043 7688141
## Dec 2014 7113767 5980360 8247174 5380371 8847164
## Jan 2015 8919230 7781688 10056772 7179509 10658950
## Feb 2015 7833393 6691730 8975055 6087371 9579415
## Mar 2015 7005154 5859386 8150921 5252853 8757454
## Apr 2015 6358706 5208848 7508565 4600150 8117263
## May 2015 6086344 4932409 7240279 4321553 7851135
## Jun 2015 7653295 6495299 8811292 5882292 9424298
## Jul 2015 8801193 7639149 9963237 7023999 10578386
## Aug 2015 9301580 8135502 10467658 7518218 11084942
## Sep 2015 8524478 7354380 9694575 6734968 10313987

```


```
## Oct 2015      6593256 5419152 7767359 4797619 8388892
## Nov 2015      5961092 4782996 7139188 4159350 7762834
## Dec 2015      7113767 5931693 8295842 5305940 8921594
```

summary.stl

```
##          Point Forecast    Lo 80      Hi 80      Lo 95      Hi 95
## Jan 2014      8919230 7832347 10006112 7256987 10581473
## Feb 2014      7833393 6742199 8924586 6164556 9502230
## Mar 2014      7005154 5909665 8100642 5329749 8680558
## Apr 2014      6358706 5258940 7458473 4676759 8040654
## May 2014      6086344 4982317 7190371 4397880 7774808
## Jun 2014      7653295 6545023 8761567 5958339 9348251
## Jul 2014      8801193 7688692 9913693 7099770 10502616
## Aug 2014      9301580 8184867 10418293 7593714 11009445
## Sep 2014      8524478 7403568 9645387 6810194 10238761
## Oct 2014      6593256 5468165 7718347 4872577 8313934
## Nov 2014      5961092 4831835 7090349 4234043 7688141
## Dec 2014      7113767 5980360 8247174 5380371 8847164
## Jan 2015      8919230 7781688 10056772 7179509 10658950
## Feb 2015      7833393 6691730 8975055 6087371 9579415
## Mar 2015      7005154 5859386 8150921 5252853 8757454
## Apr 2015      6358706 5208848 7508565 4600150 8117263
## May 2015      6086344 4932409 7240279 4321553 7851135
## Jun 2015      7653295 6495299 8811292 5882292 9424298
## Jul 2015      8801193 7639149 9963237 7023999 10578386
## Aug 2015      9301580 8135502 10467658 7518218 11084942
## Sep 2015      8524478 7354380 9694575 6734968 10313987
## Oct 2015      6593256 5419152 7767359 4797619 8388892
## Nov 2015      5961092 4782996 7139188 4159350 7762834
## Dec 2015      7113767 5931693 8295842 5305940 8921594
```

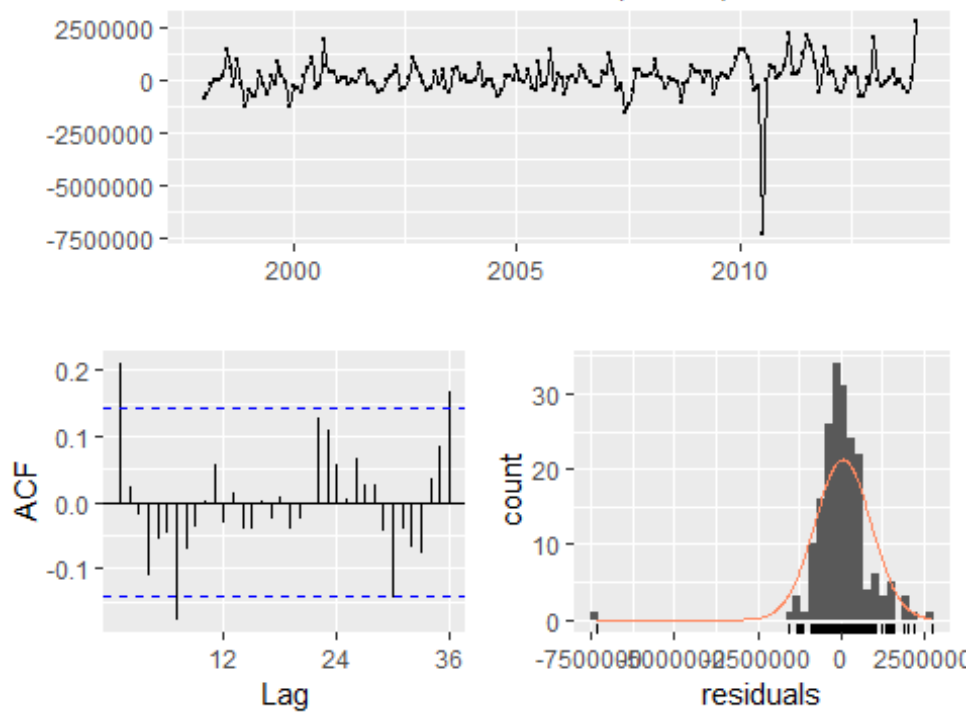
```
power.model.stl<- forecast(power.model.stl)
autoplot(power.model.stl, 12)
```

Forecasts from STL + ETS(A,N,N)



```
checkresiduals(power.model.stl)
```

Residuals from STL + ETS(A,N,N)



```
##
## Ljung-Box test
##
## data: Residuals from STL + ETS(A,N,N)
## Q* = 28.969, df = 22, p-value = 0.1457
##
## Model df: 2. Total lags used: 24
```

F. Compare Model Results/Export Data

- After comparing the RMSE of in the accuracy test, the ARIMA model will be used as the final model due to the lower RSME and better prediction capabilities. The ARIMA model also has the lowest AICc score and best score from the Box Ljung tests.

```
rmse.list <- data.frame(list(accuracy(power.model)[2], accuracy(power.model.ets)[2], accuracy(power.model.stl)[2]))
names(rmse.list)<- list('Arima', 'ETS', 'STL')
rmse.list

##      Arima      ETS      STL
## 1 595389 630869.7 843670.1
```

G. Send Results to excel

- Send to a .csv file, will manually merge into the project's consolidated file for project submission

```
write.csv(forecast.power, "Power_Forecasts_ARIMA.csv")
```


Data 624. Project 1: Part C

Team 1. Angrand, Burke, Deboch, Groysman, Karr

10/22/2019

Part C – Waterflow_Pipe1.xlsx and Waterflow_Pipe2.xlsx

Part C consists of two data sets. These are simple 2 columns sets, however they have different time stamps. Your optional assignment is to time-base sequence the data and aggregate based on hour (example of what this looks like, follows). Note for multiple recordings within an hour, take the mean. Then to test appropriate assumptions and forecast a week forward with confidence bands (80 and 95%). Add these to your existing files above – clearly labeled.

Step 1. Load Libraries

```
library(forecast)
```

```
library(ggplot2)
```

```
library(Hmisc)
```

```
library(lubridate)
```

```
library(fma)
```

```
library(readxl)
```

```
library(knitr)
```

```
library(seasonal)
```

```
library(openxlsx)
```

Step 2. Read in 2 Excel files

```
mdata1 <- read_excel("Waterflow_Pipe1.xlsx")
```

```
mdata2 <- read_excel("Waterflow_Pipe2.xlsx")
```

Step 3. Exploratory Analysis.

Let's see dimensions, top/bottom records, data types

```
dim(mdata1)
```

```
## [1] 1000    2
```

```
str(mdata1)
```

```
## Classes 'tbl_df', 'tbl' and 'data.frame':    1000 obs. of  2 variables:
```

```
## $ DateTime : POSIXct, format: "2015-10-23 00:24:06" "2015-10-23 00:40:02" ...
```

```
## $ WaterFlow: num  23.4 28 23.1 30 6 ...
```

```
kable(summary(mdata1))
```

DateTime	WaterFlow
Min. :2015-10-23 00:24:06	Min. : 1.067
1st Qu.:2015-10-25 11:21:35	1st Qu.:13.683
Median :2015-10-27 20:07:30	Median :19.880
Mean :2015-10-27 20:49:15	Mean :19.897
3rd Qu.:2015-10-30 08:24:51	3rd Qu.:26.159
Max. :2015-11-01 23:35:43	Max. :38.913

```
head(mdata1)
```

```
## # A tibble: 6 x 2
```

```
##   DateTime           WaterFlow
##   <dtm>              <dbl>
## 1 2015-10-23 00:24:06      23.4
## 2 2015-10-23 00:40:02      28.0
## 3 2015-10-23 00:53:51      23.1
## 4 2015-10-23 00:55:40      30.0
## 5 2015-10-23 01:19:17       6.00
## 6 2015-10-23 01:23:58      15.9
```

```
tail(mdata1)
```

```
## # A tibble: 6 x 2
```

```
##   DateTime           WaterFlow
##   <dtm>              <dbl>
## 1 2015-11-01 22:09:14      15.3
## 2 2015-11-01 22:09:18      26.3
## 3 2015-11-01 22:25:31      29.1
## 4 2015-11-01 23:08:20      22.8
## 5 2015-11-01 23:34:10      16.2
## 6 2015-11-01 23:35:43      21.2
```

```
dim(mdata2)
```

```
## [1] 1000      2

str(mdata2)

## Classes 'tbl_df', 'tbl' and 'data.frame':   1000 obs. of  2 variables:
## $ DateTime : POSIXct, format: "2015-10-23 01:00:00" "2015-10-23 02:00:00" ...
## $ WaterFlow: num  18.8 43.1 38 36.1 31.9 ...

summary(mdata2)

##      DateTime              WaterFlow
##  Min.      :2015-10-23 01:00:00   Min.    : 1.885
##  1st Qu.:2015-11-02 10:45:00   1st Qu.:28.140
##  Median :2015-11-12 20:30:00   Median :39.682
##  Mean   :2015-11-12 20:30:00   Mean    :39.556
##  3rd Qu.:2015-11-23 06:15:00   3rd Qu.:50.782
##  Max.    :2015-12-03 16:00:00   Max.    :78.303

head(mdata2)

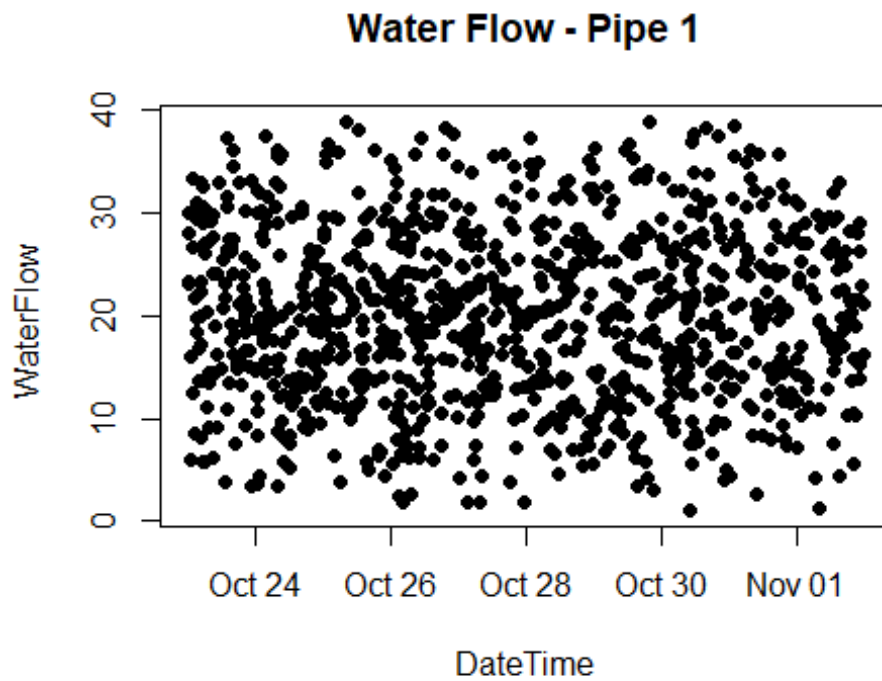
## # A tibble: 6 x 2
##   DateTime              WaterFlow
##   <dtm>                <dbl>
## 1 2015-10-23 01:00:00         18.8
## 2 2015-10-23 02:00:00         43.1
## 3 2015-10-23 03:00:00         38.0
## 4 2015-10-23 04:00:00         36.1
## 5 2015-10-23 05:00:00         31.9
## 6 2015-10-23 06:00:00         28.2

tail(mdata2)


## # A tibble: 6 x 2
##   DateTime              WaterFlow
##   <dtm>                <dbl>
## 1 2015-12-03 11:00:00         73.0
## 2 2015-12-03 12:00:00         31.5
## 3 2015-12-03 13:00:00         66.8
## 4 2015-12-03 14:00:00         42.9
## 5 2015-12-03 15:00:00         33.4
## 6 2015-12-03 16:00:00         66.7
```

Some basis scatter plots of our data.

```
plot(mdata1$DateTime, mdata1$WaterFlow, main="Water Flow - Pipe 1",
     xlab="DateTime", ylab="WaterFlow ", pch=19)
```



```
plot(mdata2$DateTime, mdata2$WaterFlow, main="Water Flow - Pipe 2",  
     xlab="DateTime", ylab="WaterFlow ", pch=19)
```



Step 4. Data Cleaning.

Let's get the first dataset in the right format. One record per an hour.

```
mdata1$WFp<-Lag(mdata1$WaterFlow,shift=1)

mdata1$DateTimep<-Lag(mdata1$DateTime)

#mydata1$myhour<-hour(mdata1$DateTime)

mdata1$mhour<-hour(mdata1$DateTime)

mdata1$mhourp<-hour(mdata1$DateTimep)

mdata1$WaterFlowN<-ifelse(mdata1$mhour!=mdata1$mhourp,(mdata1$WaterFlow+mdata1$WFp)/2,NA)

mdata1N<-mdata1[complete.cases(mdata1), ]

mdata1N$DateTimeN<-floor_date(mdata1N$DateTime,"hour")

mdata1N<-mdata1N[,c(8,7)]
```

Let's see dimensions, top and bottom records, and a plot of transformed data.

```
dim(mdata1N)

## [1] 235    2


head(mdata1N)

## # A tibble: 6 x 2
##   DateTimeN          WaterFlowN
##   <dtm>          <dbl>
## 1 2015-10-23 01:00:00      18.0
## 2 2015-10-23 02:00:00      17.1
## 3 2015-10-23 03:00:00      18.9
## 4 2015-10-23 04:00:00      15.3
## 5 2015-10-23 05:00:00      20.3
## 6 2015-10-23 06:00:00      29.6

tail(mdata1N)
```

```
## # A tibble: 6 x 2
##   DateTimeN      WaterFlowN
##   <dtm>         <dbl>
## 1 2015-11-01 18:00:00      18.3
## 2 2015-11-01 19:00:00      17.7
## 3 2015-11-01 20:00:00      19.7
## 4 2015-11-01 21:00:00      14.6
## 5 2015-11-01 22:00:00      21.0
## 6 2015-11-01 23:00:00      26.0

plot(mdata1N$DateTimeN, mdata1N$WaterFlowN, main="Water Flow - Pipe 1",
     ,
     xlab= "DateTime", ylab="WaterFlow ", pch=19)
```



Fixing missing data (4 records were missing), taking into account time zone and daylight saving time.

#4 hour difference

```
nr1<-data.frame(as.POSIXct("2015-10-27 17:00:00 -0400"),28.944308)
names(nr1)<-c("DateTimeN","WaterFlowN")
nr1

##           DateTimeN WaterFlowN
## 1 2015-10-27 17:00:00  28.94431
```

```

mdata1N <- rbind(mdata1N, nr1)

#4 hour difference
nr2<-data.frame(as.POSIXct("2015-11-01 01:00:00 -0400"),19.998079)
names(nr2)<-c("DateTimeN","WaterFlowN")

#4 hour difference
nr2<-data.frame(as.POSIXct("2015-10-28 00:00:00 -0400"),17.089225)
names(nr2)<-c("DateTimeN","WaterFlowN")

mdata1N <- rbind(mdata1N, nr2)
#5 hour difference - time change
nr3<-data.frame(as.POSIXct("2015-11-01 08:00:00 -0500"),23.474922)
names(nr3)<-c("DateTimeN","WaterFlowN")

mdata1N <- rbind(mdata1N, nr3)

mdata1N1 <- mdata1N[order(mdata1N$DateTimeN),]

```

Step 5. Converting data into time series.


Our first dataset only covers time period from 10/23/2019, 1AM to 11/1/2019 11PM, while the second dataset covers from 10/23/2019, 1AM to 12/3/2019, 4PM. We are required to forecast one week of data flow for both pipes. So, the correct way would be to predict one week from the earliest data set, or from 11/1/2019, 11PM. For that time period, we only need to forecast first dataset and for the second we have actual data.

```

ts1<-ts(mdata1N1$WaterFlowN,start=c(2015,10,23,1),freq=24)

autoplot(ts1) +
  ggtitle("Water Flow - Pipe 1") +
  xlab("Hour") +
  ylab("Water Flow")

```



```
#ts2<-ts(mdata2$WaterFlow,start=c(2015,10,23,1),freq=24*365)

#ts2

#autoplot(ts2) +
#  ggtitle("Water Flow - Pipe 2") +
#  xlab("Hour") +
#  ylab("Water Flow")

#mdataM=merge(x = mdata1N1, y = mdata2, by.x = "DateTimeN", by.y="Date
Time")

#dim(mdataM)

#dim(mdataM)

#mdataM$WaterFlowC<-mdataM$WaterFlowN+mdataM$WaterFlow

#mdataM<-mdataM[,c(1,4)]

#mdataM

#strftime(mdataM$DateTimeN,"%Y-%m-%d %H:%M:%S %z")

#ts3<-ts(mdataM$WaterFlowC,start=c(2015,10,23,2),freq=24)
```



```
#autoplot(ts3) +
# ggtitle("Water Flow - Pipe 1 and 2") +
# xlab("Hour") +
# ylab("Water Flow")
```

Step 6. Looking at seasonality and trend.


Maximum value

```
which.max(ts1)/24
```

```
## [1] 2.541667
```


The spike in water flow was on 3 day, we can see on the graph.

```
ggseasonplot(ts1)
```



No clear pattern in water use by time of day. Even though, more water seems to be used in late hours.

```
ggsubseriesplot(ts1)
```




No clear picture. But the top hours were evening and night. The lowest water use was in mornings.

```
gglagplot(ts1)
```




```
ggAcf(ts1, lag=24)
```



Again no clear pattern

Step 8. Applying decomposition.

```
ts_decomp<-decompose(ts1,type="multiplicative")  
  
autoplot(ts_decomp) +  
  ggtitle("Hourly water flow - 10 days time period - Multiplicative De  
composition") +  
  xlab("Hourly")
```




There is some type of seasonality - pattern repeats daily. But no clear trend.

Seasonally adjusted data


```
adjust_ts<-ts1/ts_decomp$seasonal
```

```
autoplot(adjust_ts) +  
  ggtitle("Hourly water flow - 10 days - Seasonally Adjusted") +  
  xlab("Hourly")
```




```
fit <- stl(ts1, t.window=13, s.window="periodic",  
  robust=TRUE)  
  
fit1<-fit %>% seasadj() %>% naive()  
  
fit1%>%autoplot() + ylab("New orders index") +  
  ggtitle("Naive forecasts of seasonally adjusted data")
```

Naive forecasts of seasonally adjusted data



```
fit2<-fit %>% forecast(method="naive")  
fit2%>%autoplot() + ylab("New orders index")
```

Forecasts from STL + Random walk



```
fcast <- stlf(ts1, method='naive')
```


Step 9. Exponential Forecasting.

Simple exponential forecast.

```
fc <- ses(ts1, h=24*7)
# Accuracy of one-step-ahead training errors
round(accuracy(fc),2)


##           ME RMSE MAE   MPE  MAPE  MASE  ACF1
## Training set  0 5.37 4.4 -8.68 25.25 0.71    0

autoplot(fc) +
  autolayer(fitted(fc), series="Fitted") +
  ylab("Water Flow - Pipes 1") + xlab("Hours")
```



```
fc <- holt(ts1, h=24*7)

fc2 <- holt(ts1, damped=TRUE, phi = 0.9, h=24*7)
autoplot(ts1) +
  autolayer(fc, series="Holt's method", PI=FALSE) +
  autolayer(fc2, series="Damped Holt's method", PI=FALSE) +
  ggtitle("Forecasts from Holt's method") + xlab("Hours") +
  ylab("Water Flow") +
  guides(colour=guide_legend(title="Forecast"))
```


```
e1 <- tsCV(ts1, ses, h=1)
e2 <- tsCV(ts1, holt, h=1)
e3 <- tsCV(ts1, holt, damped=TRUE, h=1)
mean(e1^2, na.rm=TRUE)
## [1] 29.51763
mean(e2^2, na.rm=TRUE)
## [1] 33.01115
mean(e3^2, na.rm=TRUE)
## [1] 31.76159
mean(abs(e1), na.rm=TRUE)
## [1] 4.435647
mean(abs(e2), na.rm=TRUE)
## [1] 4.668705
mean(abs(e3), na.rm=TRUE)
## [1] 4.561972
```

The simple exponential forecast appears to be the best.

```

fit1 <- hw(ts1,seasonal="additive")
fit2 <- hw(ts1,seasonal="multiplicative")
autoplot(ts1) +
  autolayer(fit1, series="HW additive forecasts", PI=FALSE) +
  autolayer(fit2, series="HW multiplicative forecasts",
    PI=FALSE) +
  xlab("Hours") +
  ggtitle("Water flow - Pipe 1") +
  guides(colour=guide_legend(title="Forecast"))

```



Step 9. Selecting Forecasting Method.

```

fit<-ets(ts1)

summary(fit)


## ETS(A,N,N)
##
## Call:
## ets(y = ts1)
##
## Smoothing parameters:
##   alpha = 1e-04
##
## Initial states:
##   l = 20.2605

```


```
##
##   sigma:  5.3919
##
##      AIC      AICc      BIC
## 2108.400 2108.502 2118.817
##
## Training set error measures:
##                                ME      RMSE      MAE      MPE      MAPE
MASE
## Training set 0.001743637 5.369158 4.397655 -8.679627 25.24513 0.713
2159
##                                ACF1
## Training set -0.004125927
```

Model selected is A and N and N

```
autoplot(fit)
```




```
cbind('Residuals' = residuals(fit),
      'Forecast errors' = residuals(fit,type='response')) %>%
  autoplot(facet=TRUE) + xlab("Hours") + ylab("")
```



```
fit1 <- fit%>%forecast(h=24*7,level=c(80,95))

fit1%>%
  autoplot() +
  ylab("Water flow - pipe 1 - 7 days forecast")
```



Step 10. Preparing the final file to be ouputed in the Excel

```
mdata1A<-mdata2[240:(239+24*7),]
```

```
fdata<-cbind(fit1,mdata1A)
```

```
#write.xlsx(fdata, "fdata.xlsx")
```